

GhostProofJob — Recruiter / Company Tier

Master Build Plan & Audit Checklist

A two-sided extension that lets verified employers post jobs directly into GPJ and discover candidates — built insert-only, candidate-first, and consent-gated, on top of the existing v96 benchmark.

Prepared for Aaliyah (Founder) & Kristina · Companion to `CLAUDE.md` and the Master Audit Checklist · Pins to the current live build — confirm exact version before Sprint R0 · Free until you're hired · No ads. No data selling. Ever.

0. How to read this & inherited rules

This document is the single source of truth for the recruiter expansion. It is a companion to `CLAUDE.md`, not a replacement — every rule in `CLAUDE.md` still applies in full. The recruiter tier is an **insert-only addition**: it must never regress a single candidate-facing feature or prior fix.

Status legend:

- BE** backend / logic / data only — zero frontend risk
- FE** frontend, but no new layout / overlay
- UI-REVIEW** touches layout, a view, an overlay, or z-index → **STOP, propose, get Aaliyah's OK before writing code**

Inherited non-negotiables (from CLAUDE.md — do not break): INSERT-ONLY · full drop-in files (never snippets) · **UI-REVIEW** gate before any visual change · boot harness must print "RAN TO COMPLETION" · `<div>` balance delta 0 · `index.html` and `GhostProofJob.html` byte-identical mirror · three version markers bumped in sync · honesty over optimism · no demo data in live views · every build ships PDFs + full files + test checklist. Vercel is case-sensitive. Drag-upload the 784 KB file — never paste.

Version alignment first. `CLAUDE.md` currently reads "Current build: v74" but the Master Audit Checklist benchmark is **v96 / build v77**. Before Sprint R0, confirm the true live build, refresh `CLAUDE.md`'s version + history, and pin all recruiter work to that number.

1. Scope & the core principle

The candidate app is a **one-sided product**: a person's data powers their own search and never leaves their control. The recruiter tier turns GPJ into a **two-sided marketplace** where candidate data becomes visible to a third party. That is a change to the core data contract the brand is named after. Every design decision below protects two invariants:

- Candidate-first.** The recruiter tier exists to *increase the job pool for job-seekers* — a second job source alongside the harvester. Recruiter features must never slow, block, or add cost to the candidate deck / Browse hot path.
- Consent-gated, never sold.** Recruiters pay for posting volume and tooling — **never for access to candidate contact info**. Contact is revealed by candidate consent, not by a paywall. The moment "pay to see the phone number" exists, "no data selling" is dead.

2. Locked design decisions (confirmed with founder — treat as founder-locked)

Area	Decision
Non-applicant visibility	Default: not surfaced. A candidate opt-in toggle ("Let verified employers find me & reach out about matching roles"), OFF by default, revocable anytime , makes them discoverable with full card. The toggle opens the channel + contact permission — it never auto-accepts a specific interview . Without it, non-applicants surface anonymized (match % + skills only) or not at all.
Contact reveal	Recruiter sees contact if: candidate applied to that job OR discovery-consent ON OR accepted a reach-out . Independent of the F-ADDR résumé "show phone" toggle (that governs the downloadable résumé's look only).
Reverse matching	Multi-level: (a) on job post/edit, (b) nightly delta batch, (c) on candidate résumé update. Local-market-scoped to the job's market (remote job → remote pool). Backend/async only. Top-N (~50) per job. Uses lightweight per-candidate match-token docs to keep reads down.
Verification (v1)	Corporate-domain email verification (block free/burner) + <code>isValidated:false</code> hidden-until-verified + manual admin spot-check for the first cohort + community-flag backstop. "Verified Employer" = domain-verified + admin-reviewed, stated honestly (not a legitimacy guarantee).
Company reputation	Two separate signals. (1) Community Ghost Reports / Vibe reviews — subjective, companies can never edit or delete . (2) Responsiveness Rate — objective, % of applicants who got any response in N days; a rejection counts (responding is the goal). Appeal system for reviews (below); appeals never delete a review.
Internal apply	Jobs carry a <code>source</code> flag. External (harvested) → routes out via <code>applyVia()</code> . Internal (recruiter-posted) → writes an internal application record; no external redirect. Internal jobs carry a "Verified Employer" green flag and must not be auto-flagged as ghost/stale.
Calendar scheduling	Deferred to a later phase (Sprint R7). v1 outreach ships with Interested / Not-interested + consented contact reveal. Full framework is roadmapped in §11 so the picture is complete.
Pricing	Free (2–3 active jobs) · Growth (~\$49–79/mo, ~10 jobs + Anti-Ghosting Badge) · Scale (~\$149–199/mo, unlimited + bulk import + priority verification) · optional à-la-carte per-job post. Numbers to validate before charging. See §8.

Review appeal system — valid-reason parameters (founder spec)

Companies file an "Appeal Review" request → routes **only** to the admin queue. Approval **does not remove** the review: it stays visible, tagged "Appealed · GPJ Approved," showing the appeal reason + admin decision, and only then is excluded from the score. Denied appeals also display ("Company appealed; GPJ upheld"). Valid reasons are a **fixed enum**:

- 1. Factually false** — a specific claim contradicted by system records (e.g. "never posted this job" when a record exists).
- 2. Not a genuine applicant** — reviewer never applied to or interacted with this company (no matching record).
- 3. Mistaken identity** — review left on the wrong company / wrong role.
- 4. Prohibited content** — harassment, hate, doxxing, personal data, profanity, spam / off-topic.
- 5. Duplicate / manipulation** — same user spamming, or coordinated brigading.

Explicitly NOT valid grounds: "the review is negative," "we disagree," "it makes us look bad." Honest-but-unflattering feedback stays and keeps counting.

3. Data model additions (all NEW — nothing today has a recruiter concept)

Collection / field	Purpose
recruiters/{uid} (or role on profiles)	Recruiter account: company name, verified website, LinkedIn, team info, working hours/timezone, role: 'recruiter' , plan , isValidated , availability grid (deferred).
companies/{companyId}	Canonical company record (may reuse the existing company/ghost namespace). Holds Responsiveness Rate, Verified-Employer status, appeal decisions.
jobs/{jobId} — new fields	source: 'internal' , ownerId , companyId , isValidated , status (open/closed/filled), market . Reuse existing job-card fields so the internal wizard maps 1:1 to the aggregator schema.
jobs/{jobId}/applications/{candidateUid}	Internal application: submitted résumé/cover-letter variant reference, status, timestamps. Recruiter (owner) reads; candidate reads own; no cross-read.
jobs/{jobId}/recommended_candidates/{uid}	Pre-computed reverse matches (score, matched/missing tokens, applied-flag). Backend-written only.
profiles/{uid} — new field	discoverable: false (the consent toggle) + match-token cache (pre-extracted title/skill tokens for cheap scoring).
notifications/{uid}/items/{id}	Candidate inbound tray: reach-out, rejection, interview updates. Opening auto-logs the response + updates the Responses bucket + Responsiveness Rate.
appeals/{id}	Review appeal requests → admin-only read/write. Stores reason enum, evidence, decision, decided-by.
schedules/* (deferred)	Availability slots, bookings, .ics payloads — Sprint R7 framework only.

Sprint R0 confirms the real schema. Claude Code reads the live index.html to document the exact resumeData / profiles fields and confirm whether the gpj_optimized tailored-variant shelf is keyed to (candidateUid, jobId) . If not, that keying is added there (small, candidate-side, insert-only) so "recruiter sees the résumé built for THIS job" works.

4. Firestore security rules — the enforcement layer BE

Current rules have **no recruiter concept**; profiles/{userId} is strictly self-or-admin. The privacy promise is *only* as real as these rules. Required additions:

- A recruiter reads jobs/{jobId}/applications/* and recommended_candidates/* **only for jobs they own** (ownerId == request.auth.uid) — never arbitrary profiles .
- A candidate creates/reads **their own** application doc; cannot read another candidate's application.
- Candidate profile fields exposed to a recruiter flow through a **curated, rules-safe projection** (match token doc / candidate card doc), never a raw profiles read. Contact fields gated by the consent conditions in \$2.
- recommended_candidates and match-token docs are **written only by the backend** (Cloud Function / Worker via Admin SDK, which bypasses rules) — allow write: if false for clients, mirroring the existing ai_sentence_cache / user_usage pattern.
- appeals — recruiter creates own; read/update admin-only (reuse isAdmin()).
- Job writes: recruiters write **only their own** jobs (today jobs is write: false for all clients — this is a scoped opening that must not let a recruiter touch harvested jobs).

5. Asset reuse map (reverse-engineer, don't rebuild)

Existing asset	Reused for
computeMatch / searchRankJobs / _resumeFieldWord (BLEND logic)	Inverted into the reverse-match engine (job → candidates), same blend of title + content.
Universal job card / openBrowseExpanded (F-CARD)	The Candidate Card mirrors the job card structure (header, summary, education, match grid).
vibe-review-modal (z-360) + rate-ex flow	Company reviews + the appeal-request entry point.
Existing modal/prompt components, brand tokens (#120F1D / #00F5A0 / #B55FE6 / #FF4D6A)	Internal job-creation wizard, reach-out / rejection editors, upgrade modals.
Admin / booster queue pattern + isAdmin()	Verification spot-check queue + appeal review queue.
Resend email + Cloudflare Worker (JWKS-verified) + user_usage caps	Outreach emails, backend match scoring, rate-limiting recruiter actions.
Community ghost-report system	Anti-junk backstop for internal jobs.
Applied/Skipped/Responses + renderStatList	Candidate notification tray + auto-response logging extend these, not replace them.

6. Master checklist — by sprint / version

BE backend/logic FE frontend, no new layout UI-REVIEW propose before coding

Stabilize before building on it. Live testing of v96 surfaced candidate-side gaps the recruiter tier would inherit (it reuses the job card, job data, company card, review flow, and AI). Sprint R-pre resolves and *re-verifies* these first, and reconciles the checklist's status flags against what's actually live, so the recruiter tier stands on a confirmed-solid base. R-pre gates R1+; non-blocking candidate items (F-TPL, F-DESK, B4) stay in the normal candidate queue.

Sprint R-pre — Stabilize & verify the v96 benchmark (candidate-side; gates R1+)

- ☐ **Checklist status reconciliation:** prove each Section-4 "done" item is actually live (F-CARD uniformity, review flow, F-RATER/F-AI/F-COVERLETTER); mark the true state — the status boxes have drifted from shipped work **BE**
- ☐ **Smart job-data gathering (root cause = harvester storage, NOT the card):** keep extracting by section (desc/req/benefits/responsibilities) and **strip boilerplate before storing** — EEO/legal statements, generic "about us" fluff, apply instructions — which shrinks docs AND sharpens matching (leaner text helps `computeMatch` / `_reqGaps`). The card already shows all it stores; the "... " ending is truncated source data **BE**
- ☐ **Sample live postings first, then set generous per-section ceilings that rarely fire** (starting targets: desc ~10k · req ~3k · benefits ~2k · responsibilities ~3k); keep display caps above them; the rare overflow tail is covered by the existing *View Full Posting* link (never store the whole raw page) **BE**
- ☐ **Backfill stale docs:** old jobs stored under small caps / empty benefits only heal on re-harvest — define a cost-aware backfill pass (ties to D1) so existing cards fill in **BE**
- ☐ **Requirements/qualifications coverage:** with fuller stored req text, `_reqGaps` sees degree/experience/cert/skill requirements it currently loses to truncation — re-verify the chip appears where it should **BE**
- ☐ **Benefits/bullets render:** confirm F-STRUCT bullets + benefits display once the data is stored (structure shipped v81; the gap is empty/truncated source data) **FE**
- ☐ **Review/vibe dedup + company-card uniformity:** keep the top vibe stars, remove the duplicate review control; company card opens in ONE format from every entry point (job card already uniform) — re-open F-CARD's uniformity piece and verify **UI-REVIEW**
- ☐ **F-ADDR address toggle** (requested, never built): show-address on/off (default off) · full-address vs city/state · layout choice; must not break export/spacing/headshot **UI-REVIEW**
- ☐ **Max Distance slider — remove now.** It can't filter on the current text/metro model (no coordinates on docs), so it's a dead control. Deferred to **F-GEO** (future version) for a proper rebuild **FE**
- ☐ **Smart-AI live verification:** confirm every section's "improve" action returns a real, section-appropriate result (Jett summary, key duties, Match-to-Job, cover letter) — test across sections, not just variation-tracking **BE**
- ☐ Benign self-test warning (Day-N vs "Unlimited" header) reconciled to a green self-test **FE**

Sprint R0 — Foundation + test scaffolding (backend/schema, zero frontend risk)

Scaffolding-first. Build the test plumbing that doesn't need UI *before* any feature code, so every later sprint has a ready test lane and we avoid build → test → refix churn. This also closes the pre-existing **F-TEST** backend-coverage gap, not just recruiter-flavored copies of existing tests.

- ☐ Confirm live build; refresh `CLAUDE.md` version + history **BE**
- ☐ Document exact `profiles` / `resumeData` schema from live code; confirm `gpj_optimized` variant keying **BE**
- ☐ Add recruiter data model + role field; no candidate-path reads changed **BE**
- ☐ Draft the full drop-in security-rules file (\$4) **BE**
- ☐ **Firestore emulator wired into CI** + `@firebase/rules-unit-testing` rules test matrix — prove a recruiter cannot read another recruiter's applicants or any raw candidate profile, a candidate cannot read others' applications, and client writes to backend-only docs are rejected **BE**
- ☐ **Fixture / seed + teardown helpers** (test company, internal jobs, matched candidate) — emulator only, **zero live-data pollution** (respects "no demo data in live views") **BE**
- ☐ **Backend match-path test skeleton** + architectural rule: author the reverse-match scorer in-repo so it is unit-testable, not only in the opaque Worker **BE**
- ☐ Stub the recruiter Playwright project in `playwright.config.js` ; document the CI secret that activates the authed-recruiter quadrant **BE**
- ☐ Multi-role state-coverage matrix (guest / authed-candidate / authed-recruiter / empty / network-fail) written before any code **BE**

Sprint R1 — Recruiter onboarding & verification

- ☐ Dual-role signup/login route (`role:'recruiter'`); candidate path byte-untouched **UI-REVIEW**
- ☐ **Recruiter auth harness** (mirrors `auth.setup.js`) — signs in a test recruiter + saves session (incl. IndexedDB); **activates the authed-recruiter quadrant** the moment the route exists **BE**
- ☐ Company profile dashboard (name, website, LinkedIn, team, hours/TZ) **UI-REVIEW**
- ☐ Corporate-domain email check + free/burner block (MX/domain lookup) **BE**
- ☐ `isValidated:false` gate — recruiter jobs hidden from candidates until verified **BE**
- ☐ Admin verification queue (reuse booster-queue pattern) **UI-REVIEW**

Sprint R2 — Job creation, import & internal apply

- ☐ Internal job-creation wizard — fields map 1:1 to the aggregator job-card schema **UI-REVIEW**
- ☐ Bulk CSV importer with field-mapping **UI-REVIEW**
- ☐ `applyVia()` branches on `source` ; internal apply writes application record + notifies recruiter **FE**
- ☐ Verified-Employer green flag; ghost-risk heuristic excludes verified internal jobs **FE**
- ☐ Free-tier job limit interceptor (upgrade modal); payment processor deferred to R8 **FE**
- ☐ Job lifecycle: close/fill → hide + notify applicants (anti-ghosting) **BE**

Sprint R3 — Reverse-match engine (backend)

- ☐ Invert match logic into a backend scorer (Worker/Cloud Function), local-market-scoped **BE**
- ☐ Match-token doc per candidate (cheap re-scoring) **BE**
- ☐ Three triggers: on job post/edit · nightly delta batch · on résumé update **BE**
- ☐ Write top-N to `recommended_candidates/` ; cost logging into D1 tracking **BE**
- ☐ Guardrail test: candidate deck/Browse read counts unchanged **BE**

Sprint R4 — Recruiter dashboard & candidate cards

- ☐ Dual-view command center (jobs list · matched candidates) **UI-REVIEW**
- ☐ Candidate card mirroring the job card + applied indicator **UI-REVIEW**
- ☐ 1-to-1 fit: match % + matching (green) / missing (grey) skills grid + pros/cons **UI-REVIEW**
- ☐ Privacy gates: contact masked per \$2 consent rules **FE**
- ☐ Context-aware doc viewer — the exact tailored résumé/cover letter for that job **FE**

Sprint R5 — Outreach & anti-ghosting

- ☐ Reach-out (interview) modal: editable, name-dynamic template UI-REVIEW
- ☐ Rejection modal: template + optional reason (shown to applicant only if filled) UI-REVIEW
- ☐ Passive candidates (never applied): "Not a Match" → silent skipped queue, no rejection button, no alert FE
- ☐ Both send via email (Resend) + internal notification BE
- ☐ Responsiveness Rate metric + Anti-Ghosting Badge (one response per applicant) BE
- ☐ Appeal request flow + admin approve/deny + visible appeal tag (never deletes review) UI-REVIEW

Sprint R6 — Candidate notifications, consent & copy sweep

- ☐ Candidate discovery-consent toggle at signup/résumé setup (OFF default, revocable) UI-REVIEW
- ☐ Candidate notification tray/dashboard (bell + list; extends Responses) UI-REVIEW
- ☐ Opening a notification auto-logs response + updates Responses + Responsiveness Rate BE
- ☐ Interview response: Interested → reveal contact / Not-interested → remove job from deck/Browse/swipe FE
- ☐ Transparency/copy sweep across site (§9) FE

Sprint R7 — Calendar scheduling (deferred phase — see §11)

- ☐ Recruiter availability grid (weekly hours, slot length, timezone) UI-REVIEW
- ☐ Candidate slot picker + propose-alternate; transaction-locked to prevent double-booking UI-REVIEW
- ☐ .ics generation (open-source, in serverless) + dual-party confirmation email BE

Sprint R8 — Monetization & billing

- ☐ Payment processor (Stripe or equivalent) — subscription + à-la-carte BE
- ☐ Plan-limit enforcement wired to real billing state FE
- ☐ Upgrade / downgrade / cancellation flows UI-REVIEW

Deferred — candidate roadmap (tracked here for completeness)

- ☐ **F-GEO — geolocation-backed distance filter.** Add lat/long to the user + job docs via a bundled US city-centroid table (free, offline, no API cost); compute haversine distance at scope time; **re-enable the Max Distance slider** once it truly filters. Bonus: sharpens recruiter local-market scoping (R3). Lands on a later candidate version, ideally riding along with harvester work BE

7. Per-build regression protocol (extends CLAUDE.md §4)

Every recruiter build runs the standard benchmark **plus** a multi-role pass, because there are now two account types sharing one file:

- Standard: JS syntax → **boot harness "RAN TO COMPLETION"** → `<div>` delta 0 → mirror identical → handler audit → dup-ID check → 3 version markers → in-app self-test.
- Multi-role state matrix** — every change verified across: guest · authed-candidate · authed-recruiter · empty-data · network-fail. The candidate quadrants must show **zero behavior change**.
- Playwright: existing candidate suite stays 100% green (regression guard) + recruiter-role specs (signed-in recruiter harness, mirroring the candidate auth.setup pattern).
- Rules layer:** the Firestore emulator + `@firebase/rules-unit-testing` matrix runs every build that touches rules or the data model — the boot harness does NOT exercise rules or the Worker, so this is the only real proof of the privacy isolation.
- Backend match layer:** the in-repo reverse-match scorer has unit tests over fixtures (local-scope, remote-scope, empty, applied-flag) — run whenever matching changes.
- Cost check: after any matching change, confirm candidate hot-path Firestore reads are unchanged; log recruiter-matching reads to D1.
- One clean verified push per version · mirror identical · markers bumped · PDFs + test checklist shipped.

8. Monetization detail

B2B revenue is fully compatible with "free until you're hired / no ads / no data selling" — that promise is to *candidates*. The guardrail: **recruiters pay for volume + tooling + reach, never for candidate contact**.

Tier	Price (validate)	Active jobs	Includes
Free	\$0	2-3	Post, matched candidates, basic dashboard. The supply loss-leader that fills your job pool.
Growth	~\$49-79/mo	~10	+ Verified Anti-Ghosting Badge, responsiveness analytics.
Scale	~\$149-199/mo	Unlimited	+ Bulk CSV import, priority verification, (later) scheduling.
À-la-carte	~\$15-25 / 30-day post	1	One-off hirers who won't subscribe — captures the long tail, still feeds the pool.

The **premium hook is the Anti-Ghosting Badge, not job count** — it's unique to GPJ and aligns paid behavior with the mission (responsive employers get rewarded and attract talent). At launch, lean generous on the free tier: internal jobs are worth more as candidate inventory than the early subscription revenue. (Pricing, tax, and the employment-law dimension of recruiter screening warrant professional advice before you charge real cards — this is a strategy recommendation, not legal or financial advice.)

9. Transparency / copy sweep FE

The launch doc and site currently promise "your data powers your job search — nothing else." That becomes inaccurate unless reworded. This is a required sweep, on-brand with the no-misleading-copy rule:

- Update the data-use promise: applying to a job, or turning on discovery, makes your profile visible to that verified employer — **still never sold, never advertised**.
- Add plain-language copy at the consent toggle explaining exactly what "yes" does and that it's revocable.
- New account-type language across signup, pricing, and About (candidate vs. recruiter).
- Honest "Verified Employer" definition wherever the badge appears.
- Recruiter-facing copy makes clear contact is consent-gated, not purchasable.

10. Open risks & honest limitations

- Cost cliff:** Blaze trial credit ends **2026-09-19**. Recruiter matching adds reads; D1 read-cost work must land before or alongside R3, or this becomes a real bill.
- Thin density at launch:** few recruiters/candidates per metro → many jobs show few/zero matches. Honest empty states, no filler.
- Verification isn't a guarantee:** domain + admin review reduces junk but can't fully stop a determined scammer or wrongly-block a legit Gmail-using small business — say so in copy; keep the community backstop.
- Matching is heuristic:** surfaced as "suggested fit," never a promise — same honesty rule as the candidate side.
- Scale is real work now scoped small:** notification tray, dashboard, and reverse-matching are genuinely new surfaces — each gated by UI-REVIEW so nothing ships speculative.

11. Calendar scheduling framework (deferred — Sprint R7, roadmapped for the full picture)

- Availability config grid** on the company dashboard: weekly open hours, slot length (15/30/60), timezone.
- Candidate scheduling console:** on an interview invite, lock an open slot or propose an alternate back to the recruiter.
- Double-booking protection:** Firestore transaction locks a slot on selection.
- Universal .ics generation** in serverless (open-source, no paid calendar API) — syncs Google / Apple / Samsung.
- Dual-party confirmation email** (Resend) with the `.ics` attachment to both sides.
- Fallback (v1 behavior):** if the recruiter hasn't set availability, the candidate only gets Interested / Not-interested; Interested reveals contact so the recruiter reaches out directly.